



BUS 243

Lecture 3: Text Representation

PROCESS

- Tokenization is a particular kind of document segmentation
 - Break up text into smaller chunks
 - Document → paragraphs → sentences → phrases → tokens
- Tokenization is the first step in an NLP pipeline
 - turns an unstructured string (text document) into a numerical data structure
- Let's see the code



TERMINOLOGIES...

- “Type,” “word,” and “token” are terms used to describe different units of text
- Type: a unique distinct word or lexeme in a language, independent of inflections
 - *The cat chased the cat*
 - Two types: “cat” and “chased”
- Word: a grammatical unit that represents a single lexical unit of meaning
 - *I like to read books*
 - Five words
- A token is an instance or occurrence of a word
 - *I like to read books.*
 - 6 tokens including a punctuation
 - Use token and word interchangeably here



ONE-HOT VECTOR?

- What's the downside?
 - Storing all those zeros, and trying to remember the order of the words in all your documents, doesn't make much sense
- You'd like to compress your document down to a single vector rather than a big table
 - Trade-off: need to give up something
- We will revisit this representation later (sentence-base analysis)
 - Essential input for CNN (Convolution Neural net)
- Let's go back to the code



BAG OF WORDS

- A common approach is to use a column vector of word counts
 - $x = [0, 1, 1, 0, 13, \dots]^T$, where x_j is the count of word j
 - The length of x is the set of possible words in the vocabulary
- x is a vector, but it is often called a **bag of words**
 - Includes only information about the count of each word
 - NOT the order in which the words appear
 - Ignore grammar, sentence boundaries, paragraphs



TOKENIZATION

- The first subtask for constructing a BOW vector is **tokenization**
 - A sequence of characters → a sequence of **word tokens**
- Note whitespace-based tokenization is not ideal
- Tokenization is typically performed using regular expressions, with modules designed to handle each cases
 - Go back to the code example



TOKEN IMPROVEMENT

- See the text for the regular expressions
- See a number of tokenizers in the code examples
- Social media researchers have found that emoticons and other forms of orthographic variation pose new challenges for tokenization, leading to the development of special purpose tokenizers to handle them
 - O'Connor, B., M. Krieger, and D. Ahn (2010). Tweetmotif: Exploratory search and topic summarization for twitter. In Proceedings of the International Conference on Web and Social Media (ICWSM), pp. 384–385.



TOKENIZATION IS HARD?

- Tokenization is a language-specific problem
 - Each language poses unique challenges
 - Chinese does not include spaces between words, nor any other consistent orthographic markers of word boundaries
 - German does not include whitespace in compound nouns
- Social media raises similar problems for English and other languages
 - #TrueLoveInFourWords
 - Decomposition analysis (Brun and Roux, 2014)



EXTENDING YOUR VOCABULARY WITH N-GRAMS

- Now we consider a sequence of words
 - Ice cream
 - Boston Red Sox
- N-gram is simply a sequence of n words
 - N-gram could denote characters, but focus on words now
- We have tokenized sentences using 1-gram only thus far
- Using 2-gram or 3-gram words means adding more tokens in the vocabulary
 - Not difficult to add (see the codes)
 - N-gram tokens are pretty rare → need some ways to handle them properly



WHAT IF ONLY USE 1-GRAM TOKENS?

- What is the problem of rare 2-grams when we add them in the vocab?
 - Again, they are so rare. Why is this a problem?
- If use 1-gram tokens only, the stop words are usually counted the most
 - The, a, an, ...
- If they are removed:
 - Mark **reported** to the **CEO**
 - Suzanne **reported** as the **CEO** to the board
 - Lack of information about the professional hierarchy
- If not, the length of vocabulary would be the problem
- Let's dig this issue a little bit deeper



TEXT NORMALIZATION

- After splitting the text into tokens, the next question:
 - Which tokens are really distinct
- Complete elimination of case distinction will result in a smaller vocab
 - Necessary to distinguish *great*, *Great* and *GREAT*?
 - How about *apple* and *Apple*?
- Text normalization refers to string transformation that remove distinctions that are irrelevant to downstream applications
 - Also include standardization of numbers (1,000 → 1000) or dates
 - Social media (e.g. coooooooooool)



INFLECTIONS MATTER?

- A more extreme form of normalization is to eliminate inflectional affixes (e.g. -ed and -s suffixes)
 - Whale, whales, whaling all refer to the same underlying concept
- A stemmer is a program for eliminating affixes
 - Apply a series of regular expression substitutions
 - Character-based stemming algorithms are necessarily approximate

Original	Th e	Willia ms	sister s	are	Leavin g	This	Tennis	centr e
Porter stemmer	the	william	sister	are	leav	thi	tenni	centr
Lancaster stemmer	the	william	sist	ar	leav	thi	ten	cent
WordNet lemmatizer	The	William s	sister	are	leaving	this	tennis	centre



INFLECTIONS MATTER?

- **Lemmatizers** are systems that identify the underlying lemma of a given wordform
 - Geese → Goose
- Generalization would matter
 - Even inaccurate stemming can improve bag-of-words classification
 - merging related strings and thereby reducing the vocabulary size
 - However, need to avoid the over-generalization errors
- Both stemming and lemmatization are language-specific
 - English stemmer or lemmatizer is of little use on a text in another language




NORMALIZATION IS KIND OF SMOOTHING

- The value of normalization depends on the data and the task
 - Normalization reduces the size of the feature space
 - Can help in generalization
- There is always the risk of merging away meaningful distinctions
- In supervised machine learning, regularization and smoothing can play a similar role to normalization
 - Mitigate overfitting to rare (language-specific) features
- In unsupervised learning, such as topic modeling, normalization is even more critical



HOW MANY WORDS?

- Limiting the size of the feature vector reduces the memory and increases the speed of prediction
- Normalization can help to play this role, but a more direct approach is simply to limit the vocabulary to the N most frequent words in the dataset



STOPWORDS

- Another way to reduce the size is to eliminate stopwords (the, to,...)
 - Typically done by creating a stoplist (nltk stopwords) and ignoring all terms that match the list
- However, seemingly inconsequential words can offer surprising insights about the author or nature of the text (Biber 1991)
- High-frequency words are unlikely to cause overfitting in discriminative classifiers
- As with normalization, stopword filtering is more important for unsupervised problems, such as term-based document retrieval

